

Containerization for R&D Applications (Part 1!)

Daniel Trahan

Daniel.Trahan@colorado.edu

Please sign in:

<https://tinyurl.com/y4aks9zs>

Slides and exercises available for download at:

https://github.com/ResearchComputing/CONTAINER_TUTORIAL_FALL_2020

Outline

Part 1: Container fundamentals and Docker (11/12/20)

- Introduction to containers
- Docker commands and options
- Hands-on: Running Docker containers on your personal Machine
- Hands-on: Building Docker images
- Hands-on: Practical Application

Part 2: Containers for HPC with Singularity (11/19/20)

- Singularity commands and options
- Hands-on: Running containers
- Building containers
- Special cases: Running containers for MPI and GPU jobs

Introduction to Containers



What is a container?

A container is a portable environment that packages some or all the following: an operating system, software, libraries, compilers, data and workflows. Containers enable:

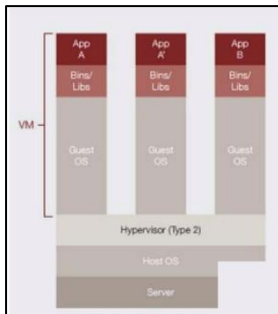
- Mobility of Compute
- Reproducibility (software and data)
- User Freedom



Virtualization (1)

Hardware virtualization (not used by containers!)

- Can run many OS's on same hardware (machine)
- E.g., VirtualBox, VMWare



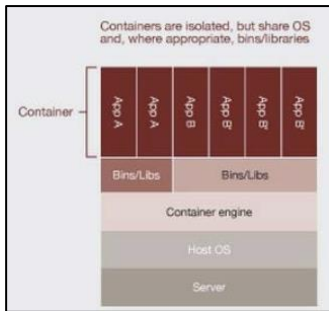
Material courtesy: M. Cuma, U. Utah

Virtualization (2)

OS-level virtualization (used by containers!)



- Can run many isolated OS instances (guests) under a server OS (host)
- Also called containers
- E.g., Docker, Singularity

Best of both worlds: isolated environment that user wants, but can leverage host OS resources (network, I/O partitions, etc.)



Material courtesy: M. Cuma, U. Utah

Containerization software

- Docker 
 - Part 1 focus of today's tutorial
 - Well established – largest user base
 - Has Docker Hub for container sharing
 - Problematic with HPC
- Singularity 
 - Part 2 focus of next week tutorial
 - Designed for HPC
- Charliecloud; Shifter
 - Designed for HPC
 - Based on Docker
 - Less user-friendly

Installing Docker

- Docker Community Edition
 - Windows: Windows 10 Professional or Enterprise
 - Mac: OS X El Capitan 10.11 or later
 - Linux
- Docker toolbox
 - Legacy solution for Windows and Mac for versions that do not meet the version requirements.
 - Utilizes the Virtual Box hypervisor for virtualization
 - For this tutorial, Windows users regardless of version will use Docker toolbox



Why Docker?

- Probably the most popular containerization software
- Offers a variety of prebuilt images including:
 - Python
 - Perl
 - NodeJS
 - Ubuntu
- Very well documented with a large community creating and supporting docker images.
- DockerHub



Docker Nuts and Bolts

- Docker runs on a concept of images and containers.
 - **Images**: Saved snapshots of a container environment.
 - Made from Dockerfile or pulled from Docker Hub
 - Stored in the Docker cache on your disk
 - **Containers**: Instances of images that are generated by Docker when an image is 'run'
 - Instance of image running in memory
 - Ephemeral Instances that cannot be continued
 - Can be run interactively

Docker Commands

- Docker Commands are usually in the form of:
`docker <sub-command> <flags> <target/command>`
- Examples:
`docker run -it myimage`
`docker container ls`
`docker image prune`

Running Docker Containers

- Run a docker image as a container:

```
docker run <image-name>
```

- Run a docker image interactively:

```
docker run -it <image-name>
```

- If an image is not on the system, then Docker will search Dockerhub to see if the image exists.
- Specify commands after your image to execute specific software in your container.

```
docker run <image-name> <program>
```

Containerized Hello World

- Let's start with something simple:
 - Docker "Hello, World!"
 - Relatively small image
 - No dependencies
 - Built as a general test case
- Command:

```
docker run hello-world
```

Docker Image and Container Commands

Image Commands	
<code>docker image ls</code>	List docker images stored in cache:
<code>docker image rm <image></code> <code>docker rmi <image></code>	Remove (an) image(s):
<code>docker image prune</code>	Remove unused images
Container Commands	
<code>docker container ls</code>	List docker containers currently running:
<code>docker container rm <container></code> <code>docker rm <container></code>	Remove (an) container(s):
<code>docker container prune</code>	Remove all stopped containers

Docker Utility Commands

Commands	
<code>docker info</code>	Shows Docker system-wide information
<code>docker inspect <docker-object></code>	Shows low-level information about an object
<code>docker config <sub-command></code>	Manage docker configurations
<code>docker stats <container></code>	Shows container resource usage
<code>docker top <container></code>	Shows running processes of a container
<code>docker version</code>	Shows docker version information

- More details and commands can be found [on the docker documentation page](#)

Demo 1: Running Containerized Python



Demo 1: Python

- Running the Python docker container will pull Python from Docker Hub:

```
docker run python:3.7.2-slim
```

- ...did it work?
- Run your python image interactively:

```
docker run -it python:3.7.2-slim
```

- This puts us into a python interpreter, where you can run python code containerized in its own environment.

Building Docker Containers

- To build a docker container, we need a set of instructions Docker can use to set up the environment.
 - Dockerfile
- Once we set up our dockerfile we can use the command
`docker build -t <image-name> .`
- Then we can run the image with our `docker run` command
`docker run <image-name>`

Demo 2: Building an Ubuntu Container



Demo 2: Setup

Dockerfiles and test files are provided for this workshop. We can pull the files from a github repository as such:

1. Navigate to your home directory

```
cd ~
```

2. Clone the repository

```
git clone
```

```
https://github.com/ResearchComputing/Container\_tutorial\_Spring\_2020
```

Demo 2: Ubuntu w/ GCC

- For this first example we will be building a custom Ubuntu image that will provide a location to run the GNU Compiler Collection.
- Dockerfile provided
- Need to build:
 1. Navigate to the directory:
`cd ~/Container_tutorial_Spring_2020/dockerdemo/ubuntu-gcc`
 2. Build the docker image with:
`docker build -t happy-gcc .`
 3. Run the docker image as a container:
`docker run -it happy-gcc`

Editing Docker Images

- Suppose you have an existing docker image and want to make changes...
 - Rebuild Dockerfile!
 - Usually a bit cumbersome
 - No Dockerfile?
- Use docker commit!
`docker run -it <image-name> bash` #or any shell...
- Then commit it to the image
`docker commit <image-id> <image>`

Mounting and Accessing files

- So now that we have a working container, how can we access the test files we downloaded?
 - Mounting directories: Bind Mount
 - Allows the docker container to access files on the host OS
 - Choose host's **source directory**, files in the directory will be moved to the container's **target directory**
 - **Source Directory**: Directory on the host system.
Never within a container.
 - **Target Directory**: Directory in the Docker Container.
Never on the host system.
 - A flag set within the **docker run** command:

```
docker run --mount type=bind,source=<source>,target=<target> <image>
```

Mounting and Accessing files

- Mounting directories: Volume Mount
 - Same concept, but volumes are stored within docker cache.
 - Create Docker volumes in your terminal and link your volume directory
 - Similarly linked through the `docker run` command.

```
docker run --mount type=volume,source=<volume>,target=<target> <image>
```


Demo 2 (Cont.): Mounting

- Returning to our demo, can we give our container access to our test files?
- Let's use a bind mount!
- In the directory where our Dockerfile lives... use this command (all on one line):

```
docker run -it --mount  
type=bind,source=$(pwd)/source,target=/target happy-gcc
```

- We can `cd /target` and run our test files!
- Command:

```
gcc hello.c -o hello.exe  
./hello.exe
```

Demo 3: NCL container

- For this next example we will building a Docker image that will run the NCAR Command Language (NCL).
- Dockerfile provided
- Same process as before:
 1. Navigate to the Dockerfile found at:
[~/Container_tutorial_Spring_2020/dockerdemo/ncl](#)
 2. Build your docker file as an image titled "bright-ncl"
 3. Run your docker image as a container
- Can we test a sample script?

Dockerhub

- The place where containers live!
- Dockerhub is a Docker hosted library of public and private Docker images.
 - Free and unlimited public images
 - 1 free private repository
- Great for hosting images for fellow researchers
- Commands similar to git

Dockerhub Commands

- Download and upload docker images with ease.
 - `docker run <image>`
 - `docker pull <image>`
- Uploading a little more complicated...
 - Sign in with:
`docker login`
 - List docker images with:
`docker image ls`
 - Tag your image:
`docker tag <image-id> <your-username>/<image-name>:<tag>`
 - Push!
`docker push <your-username>/<image-name>`

Docker-Compose

- Utility that can create and install docker images.
- Builds docker images based on a docker-compose.yml file.
 - YAML: *YAML Ain't Markup Language*
 - Data serialization language
 - Describes containers you wish to build with what features.
- Not a docker command!

Docker-Compose Commands

- Build all containers in YAML file:
`docker-compose build`
- Build and run all containers in YAML file:
`docker-compose up`
- List all containers in YAML file:
`docker-compose images`
- Run a one-off command from a container:
`docker-compose run <container-name> <command>`

Thank you!

Please fill out the survey:

<http://tinyurl.com/curc-survey18>

Contact information:

Daniel.Trahan@Colorado.edu

Additional learning resources:

Slides and Examples from this course:

https://github.com/ResearchComputing/CONTAINER_TUTORIAL_FALL_2020

Web resources:

<https://training.play-with-docker.com> (docker online training materials)

<https://hub.docker.com> (Docker Hub)